

SAIR

Spatial AI & Robotics Lab

CSE 473/573

L8: FEATURE


Chen Wang

Spatial AI & Robotics Lab

Department of Computer Science and Engineering



University at Buffalo The State University of New York



Content

- Feature
 - Extraction
 - Local features, Pyramids for invariant feature detection
 - Invariant descriptors and matching
 - Harris Detection, SIFT.
 - Matching
 - Precision, Recall, F1, ROC
 - SURF, Integral Images

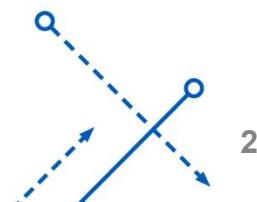


Image matching



by [Diva Sian](#)



by [swashford](#)

Harder case




by [Diva Sian](#)




by [scqbt](#)


Harder still?



NASA Mars Rover images



Answer below (look for tiny colored squares...)




NASA Mars Rover images with SIFT feature matches
Figure by Noah Snavely

Local features and alignment




- We need to match (align) images
- Global methods sensitive to occlusion, lighting, parallax.
- Look for local features that match well.
 - How would you do it by eye?




Local features and alignment

- 1. Detect feature points in both images.



Local features and alignment

- 1. Detect feature points in both images.
- 2. Find corresponding pairs.




Local features and alignment

- 1. Detect feature points in both images.
- 2. Find corresponding pairs.
- 3. Use these pairs to align images



Darya Frolova and Denis Simakov



10

Local features and alignment

- Problem 1:
 - Detect the *same* points *independently* in both images




no chance to match!

We need a repeatable detector






Local features and alignment

- Problem 2:
 - For each point correctly recognize the corresponding one




We need a reliable and distinctive descriptor

Geometric transformations



Photometric transformations



Modelled as a linear transformation:
scaling + offset

Other challenges: Noise, Blur, Compression, Artifacts, etc.

Invariant local features

Designed to be invariant to common geometric and photometric transformations.

Basic steps:

- 1) Detect **distinctive interest points**
- 2) Extract **invariant descriptors**





Figure: David Lowe

Main questions

- Where will the interest points come from?
- What are salient features that we'll *detect* in multiple views?
- How to *describe* a local region?
- How to establish *correspondences*, i.e., compute matches?



Feature Vectors – Image Templates

- How do we do image template matching (correlation)?
 - Given an image and a template, pass the template over the image
 - Find the max response over the image
- What are the challenges using image templates?
 - Must match at all locations in image.
 - Must match each template separately
 - Does not generalize
 - Scale dependent, rotation sensitive, etc..






Image Templates – Improvements*


- Matching all locations?
 - Change the Stride of matching
- Match templates separately?
 - Can do it in parallel, but same computation.
- Other Efficiency Issues
 - Extract and match other features
 - Smaller feature vector
 - Efficient Image Computation
 - Integral Images
- Generalization?
 - Rescale and rotate the images and templates

Local features: Main components

1. Detection: Identify the **interest points**.
2. Description: Extract vector **feature descriptor** surrounding each interest point.
3. Matching: Determine **correspondence** between descriptors in two views.




Detection: Corners




- Key property: in the region around a corner, image **gradient has two or more dominant directions**.
- Corners are **repeatable and distinctive**.

Corner Detection: Basic Idea


- Recognize the point by looking through a window.
- Shifting a window in *any direction* should give a *large change in intensity*.



“flat” region:
no change in
all directions



“edge”:
no change
along the edge
direction



“corner”:
significant
change in all
directions

Harris Detector

- Consider an image patch and shift it for $[\Delta x, \Delta y]$.
- *Sum of squared differences* (SSD) of two patches:

$$f(\Delta x, \Delta y) = \sum_{(x_k, y_k) \in W} (I(x_k, y_k) - I(x_k + \Delta x, y_k + \Delta y))^2$$


- $I(x + \Delta x, y + \Delta y)$ can be approximated by Taylor expansion. Let I_x, I_y be the partial derivative of I .

$$I(x + \Delta x, y + \Delta y) \approx I(x, y) + I_x(x, y)\Delta x + I_y(x, y)\Delta y$$

- This produces $f(\Delta x, \Delta y) \approx \sum_{(x, y) \in W} (I_x(x, y)\Delta x + I_y(x, y)\Delta y)^2$

$$f(\Delta x, \Delta y) \approx (\Delta x \quad \Delta y) M \begin{pmatrix} \Delta x \\ \Delta y \end{pmatrix}$$

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$



Harris Detector

- The SSD becomes:

$$f(\Delta x, \Delta y) \approx (\Delta x \quad \Delta y) M \begin{pmatrix} \Delta x \\ \Delta y \end{pmatrix}$$

where M is a 2×2 matrix with image derivatives:


$$\text{Structure Tensor } M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$

Sum over the patch



Gradient with respect to x,
times gradient with respect to y

- We sometimes add window weights to the SSD:


$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$




Harris Detector



I_x





I_y



Harris

What does this matrix reveal?

- First, consider an axis-aligned corner:




What does this matrix reveal?

- First, consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

- Dominant gradient directions align with x or y axis
- If either λ is close to 0, then this is not a corner, so look for locations where both are large.
- What if we have a corner that is not aligned with the image axes?



General Case


Since M is symmetric, we can use eigen-value decomposition

$$M = Q^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} Q$$

We can visualize M as an ellipse with axis lengths determined by the eigenvalues and orientation determined by Q matrix.


Q is orthonormal (rotation matrix)

$$Q^{-1} = Q^T.$$



Interpreting the eigenvalues

- Classification of image points with eigenvalues of M :



Harris Response Function


- The smallest eigenvalue of M:

$$R = \det(M) - \alpha \operatorname{trace}(M)^2$$

$$= \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$


- where α is an empirically determined constant.

$$\alpha \in [0.04 \text{ to } 0.06]$$



Harris corner detector: Algorithm

1. Color to grayscale
2. Calculate spatial derivative
3. Compute structure tensor for all windows
4. Harris response calculation
 - Find points with large corner response: ($R > \text{threshold}$)
5. Non-maximum suppression
 - Take the points of local maxima of R




Harris Corner Detector




Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

Harris Corner detector




Corner response R

Harris Corner detector



Find points with large corner response: $R > \text{threshold}$

Harris Detector: Workflow



Take only the points of local maxima of R

Harris Corner



Image with Harris Corner Overlay.


Harris Corner



Original Image

Harris Detector: Properties

- Rotation invariance




Ellipse rotates but its shape (eigenvalues) remains the same.

Corner response R is invariant to image rotation

Harris Detector: Properties


- Not invariant to image scale



All points will be
classified as **edges**


Corner !

- How can we detect **scale invariant** interest points?
 - **Image Pyramids!**







Invariance and Equivariance

- We want corner locations to be **invariant to photometric transformations** and **equivariant to geometric transformations**
 - **Invariance:** same corners are detected with photometric transformations, such as histogram equalization.
 - **Equivariance (Covariance):** Corners are detected **in corresponding locations** in geometrically transformed image.






Scale Invariance

- Multi-scale approach (Image Pyramids)




Scale Invariance

- Multi-scale approach (Image Pyramids)








Scale Invariance

- Multi-scale approach (Image Pyramids)







Scale Invariance

- Multi-scale approach (Image Pyramids)



Scale Invariance


- Extract patch from each image individually



Local descriptors


- We know how to detect points
- Next question:

How to *describe* them for matching?



Point descriptor should be:

1. Invariant
2. Distinctive





Local descriptors

- Simplest: list of intensities within a patch.
- What is this going to be invariant to?
 - Translation

Write regions as vectors

$$A \rightarrow \mathbf{a}, \quad B \rightarrow \mathbf{b}$$







Feature descriptors

- Disadvantage of patches as descriptors:
 - Small shifts / brightness affect matching a lot.
 - What is invariant to pixel shuffle and brightness?




- Solution: histogram of oriented gradients (HOG).



Feature descriptors: SIFT

- Scale Invariant Feature Transform
 - Divide patch into 4×4 sub-patches: 16 cells
 - Compute histogram of gradient orientations (8 reference angles) for all pixels in each sub-patch
 - Resulting descriptor: $16 \times 8 = 128$ dimensions




*	↖	*	↑
*	↖	↗	↘
*	↙	↖	↗
*	↙	*	↖


Rotation Invariant Descriptors

Dominant direction of gradient for the image patch

- Find local orientation



- Rotate patch according to this angle
 - This puts the patches into a canonical orientation.



Rotation Invariant Descriptors




Image from Matthew Brown


Feature descriptors: SIFT

- Extraordinarily robust matching technique
 - Can handle changes in viewpoint
 - Up to about 60 degree out of plane rotation
 - Can handle significant changes in illumination
 - Sometimes even day vs. night
 - Fast and efficient—can run in real time.



Working with SIFT descriptors


- One image yields:
 - n 128-dimensional descriptors:
 - each one is a histogram of the gradient orientations within a patch
 - [$n \times 128$ matrix]
 - n scale parameters specifying the size of each patch
 - [$n \times 1$ vector]
 - n orientation parameters specifying the angle of the patch
 - [$n \times 1$ vector]
 - n 2d points giving positions of the patches
 - [$n \times 2$ matrix]



Feature matching

We know how to detect **and describe** good points



Next question: **How to match them?**



Feature distance

How to define difference of two features f_1, f_2 ?



- Simple approach is $\text{SSD}(f_1, f_2)$
 - Sum of Square Differences between two descriptors
 - can give good scores to ambiguous (bad) matches



Feature distance

How to define difference between two features f_1 , f_2 ?


- Better approach:
 - ratio distance = $SSD(f_1, f_2) / SSD(f_1, f_2')$.
 - f_2 is best SSD match to f_1 in I_2 , while f_2' is 2nd best.
 - gives small values for ambiguous matches.



Feature Matching Summary

Given a feature in I_1 , how to find the best match in I_2 ?

1. Define distance function that compares two descriptors
2. Test all the features in I_2 , find the one with min difference
 - Simple approach is $SSD(f_1, f_2)$
 - sum of square differences between two descriptors
 - can give good scores to very ambiguous (bad) matches
 - Better approach: ratio distance = $SSD(f_1, f_2) / SSD(f_1, f'_2)$
 - f_2 is best SSD match to f_1 in I_2
 - f'_2 is 2nd best SSD match to f_1 in I_2
 - gives small values for ambiguous matches



Other Distance Measures

- Sum of Squared differences (SSD)


$$\text{distance}(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

- Correlation

$$\rho_{X,Y} = \text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{\mathbb{E}[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$


- Cosine Similarity

$$\text{cosine similarity} = S_C(\mathbf{A}, \mathbf{B}) := \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$



Evaluation of matches

How can we measure the performance of a feature matcher?




feature distance

True/false positives

The distance threshold affects performance

- True positives = # of detected correct matches.
- False positives = # of detected incorrect matches.




Precision, Recall, F1

$$\text{Precision} = \frac{\text{True Positive}}{\text{Predicted Positive}}$$

$$\text{Recall} = \frac{\text{True Positive}}{\text{Actual Positive}}$$

$$F_1 = \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



How many retrieved items are relevant?


$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are retrieved?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{black}}$$

Evaluation of matches



- How can we measure the performance of a feature matcher?



Evaluation of matches


ROC Curves (Receiver Operator Characteristic)

- Generated by counting # correct/incorrect matches, for different thresholds.
- Want to maximize **area under the curve** (AUC)
- Useful for comparing different feature matching methods.



Other Local Descriptors: SURF

- Speeded up robust features (SURF)
 - Fast approximation of SIFT, 6 times faster.
- Accelerated by 2D filters (Harris) & integral images.




Integral Image: Motivation

- How can we make things faster?
 - Many of the filters we have defined are made of rectangles or combinations of rectangles!




Integral Image

- A transformed image where every pixel is the sum of all pixels **above** and to the **left** of original image.




Integral image


- A quick and effective way of calculating the sum of values (pixel values)



Integral images




input image



integral image


- What's the sum of pixels in the blue rectangle?

Integral images




integral image

Integral images




integral image

Integral images




integral image

Integral images



integral image


Integral images




integral image

Integral images


- What's the sum of pixels in the rectangle?




A



B



C




D

Computing sum within a rectangle


- Let A,B,C,D be the values of the integral image at the corners of a rectangle
- Then the sum of original image values within the rectangle can be computed as:

$$\text{sum} = D - B - C + A$$

- **Only 3 additions are required for any size of rectangle!**





Integral Image Example




Why SURF is Fast?

- Second-order Gaussian derivatives and can be approximated at a very low computational cost using integral images




Why SURF is Fast?


- Instead of iteratively reducing the image size (left), the use of integral images allows the up-scaling of the filter at constant cost (right).
- The scale space is analyzed by **up-scaling the filter size**, rather than iteratively reducing the image size.
 - For each new octave, the filter size increase and sampling intervals are doubled simultaneously for the extraction of the interest points(σ), allowing the up-scaling of the filter at constant cost.



Other Types of Detector: Blob



- Blob detection
 - A Blob is a group of connected pixels sharing some common property, e.g., grayscale value.
 - Blob detection aims to identify these regions.



Important Concepts

- Key-point detection
 - Corners
 - Repeatable and Distinctive
 - Harris
- Descriptors
 - Robust and selective
 - Histograms of gradient orientation
 - SIFT, SURF
- Evaluation
 - Precision, Recall, F1, ROC

